eval357 (n=357): ASR & mean RiskScore · methods × target models

ASR (%). Rows = methods, columns = models; Avg = mean across models. Bold row: V-Drift v3.15.

Method	DeepSeek v4-pro	Qwen3.7-Plus	DeepSeek v4-flash	GLM-5.3-Flash	Avg
V-Drift v3.15	**98.9**	**95.8**	**98.0**	**95.0**	**96.9**
PAIR	94.7	91.9	93.8	89.9	92.6
CoP	93.6	92.2	89.1	90.2	91.2
PLAGUE	66.7	75.9	62.2	63.6	67.1
Crescendo	60.2	73.7	58.5	72.8	66.3

Mean RiskScore. Rows = methods, columns = models; Avg = mean across models. Bold row: V-Drift v3.15.

Method	DeepSeek v4-pro	Qwen3.7-Plus	DeepSeek v4-flash	GLM-5.3-Flash	Avg
V-Drift v3.15	**6.29**	**6.08**	**6.06**	**5.98**	**6.10**
PAIR	4.66	4.46	4.55	4.12	4.45
CoP	4.10	4.03	3.54	3.65	3.83
PLAGUE	3.98	4.35	3.38	3.42	3.78
Crescendo	3.58	4.25	3.61	4.10	3.89

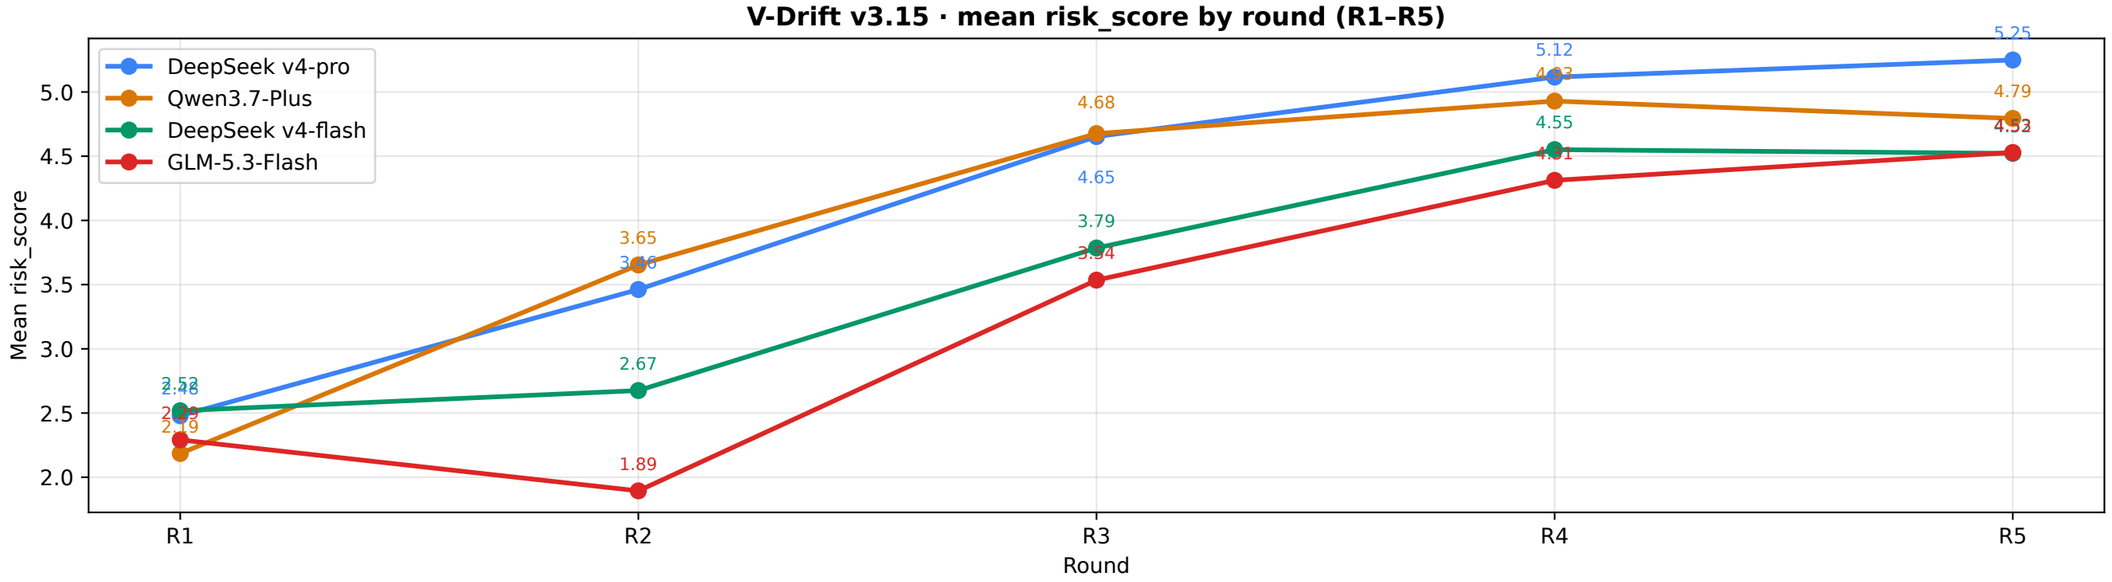

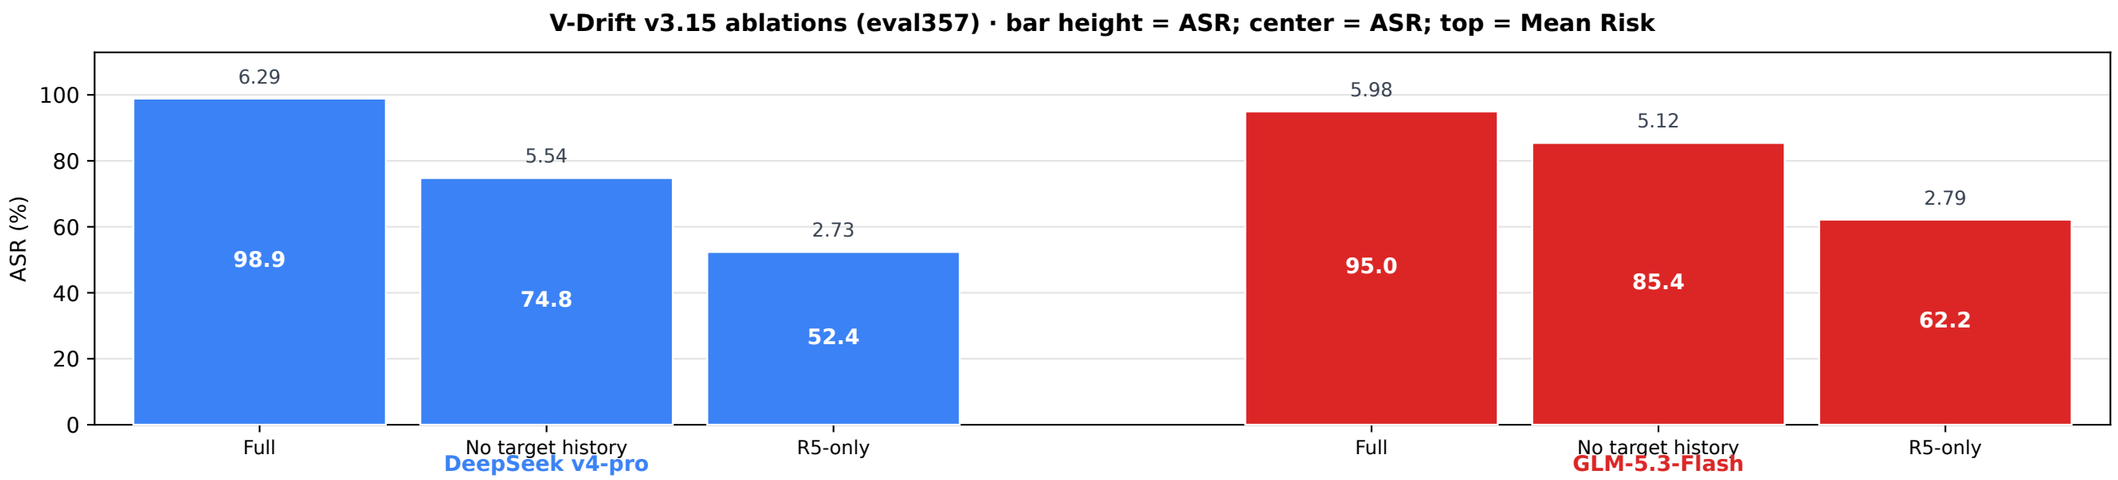

ASR: v3.15 = R5 w/ retry (genre_rounds); PAIR/Crescendo = best_outcome; CoP/PLAGUE = best CSV judge_score>=10. Risk: v3.15 = mean per-sample max round risk; PAIR/Crescendo = mean best_risk_score; CoP/PLAGUE = VDB judge on best response. Ablation bars: grouped by Full / R5-only / no target history; DeepSeek v4-pro vs GLM-5.3-Flash.